

Q1-Python

August 15, 2024

1 NAFLD detection using RNA-Seq data

Content creators: Mahdi Anvari, Sadegh Rizi

University of Tehran, Department of Biotechnology

2 Q1 - Python section

```
[1]: # Import needed libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.metrics import precision_score, recall_score
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder
from sklearn.linear_model import LassoCV
from sklearn.decomposition import PCA
from sklearn.manifold import TSNE
import warnings
warnings.filterwarnings("ignore")
from sklearn.linear_model import LogisticRegression
from sklearn.svm import SVC
from sklearn.ensemble import RandomForestClassifier
from sklearn.neural_network import MLPClassifier
```

```
[2]: # Load Data
normal_counts = pd.read_csv('Normal.counts.voom.csv')
meta_data = pd.read_csv('meta_data.csv')
```

```
[3]: normal_counts.head()
```

```
[3]:
```

	gene	DLDR_0001	DLDR_0002	DLDR_0003	DLDR_0004	DLDR_0005	\
0	ENSG000000000003	5.965571	5.741587	5.996891	5.551919	6.430237	
1	ENSG000000000005	1.612375	2.147793	0.418542	0.702492	1.215978	
2	ENSG0000000000419	4.133821	4.120969	4.086129	4.116240	4.393797	
3	ENSG0000000000457	4.111056	3.922234	3.964871	3.978350	4.018235	
4	ENSG0000000000460	4.150662	3.732756	3.634637	3.853979	3.614220	

	DLDR_0006	DLDR_0007	DLDR_0008	DLDR_0009	...	DLDR_0183	DLDR_0184	\
0	6.234619	6.071503	6.441882	5.752712	...	6.304802	6.576246	
1	0.920810	0.458163	0.927224	1.089389	...	-0.031596	-1.091275	
2	4.390909	4.148242	4.554655	4.203819	...	4.176599	4.244459	
3	3.864521	4.263119	3.896271	4.139546	...	4.378342	4.453838	
4	3.500857	4.000565	4.016287	3.904500	...	2.974209	3.720038	

	DLDR_0185	DLDR_0186	DLDR_0187	DLDR_0188	DLDR_0189	DLDR_0190	\
0	6.735760	6.344234	6.608924	6.480745	6.360397	6.367705	
1	-0.942637	-0.026585	-0.757399	-1.083676	0.886550	-0.902201	
2	4.342765	4.179319	4.274450	4.361634	4.093280	4.148010	
3	4.685598	4.438796	4.042577	4.313540	4.205119	4.506058	
4	4.640011	3.814717	2.126408	3.120196	3.336802	3.982071	

	DLDR_0191	DLDR_0192
0	6.604050	6.514539
1	-0.865036	-1.588749
2	4.351489	3.859711
3	4.072137	4.341988
4	3.143138	2.741172

[5 rows x 193 columns]

```
[4]: normal_counts.shape
```

```
[4]: (17396, 193)
```

there are 192 samples with 17396 features

```
[5]: meta_data.head()
```

```
[5]: Patient_ID    SEX    BMI_surg  Age    Run    Diabet  \
0  DLDR_0001  Female  35.214555  55  SRR8378590  Non Diabetic
1  DLDR_0002  Female  39.421748  47  SRR8378589    Diabetic
2  DLDR_0003   Male  48.758108  46  SRR8378432  Non Diabetic
3  DLDR_0004  Female  41.822607  36  SRR8378431  Non Diabetic
4  DLDR_0005  Female  53.582192  54  SRR8378434  Non Diabetic
```

	Simplified_class
0	Normal
1	Normal
2	Normal
3	Normal
4	Normal

```
[6]: meta_data.shape
```

[6]: (192, 7)

```
[7]: print("Number of Normal samples: ", meta_data['Simplified_class'].to_list().
      ↪count("Normal"))
      print("Number of Non_advanced_Fibrosis samples: ",
      ↪meta_data['Simplified_class'].to_list().count("Non_advanced_Fibrosis"))
      print("Number of Advanced_fibrosis: ", meta_data['Simplified_class'].to_list().
      ↪count("Advanced_fibrosis"))
```

Number of Normal samples: 74
Number of Non_advanced_Fibrosis samples: 53
Number of Advanced_fibrosis: 65

Let's split our data and labels to train and test sets

```
[9]: X_train, X_test, y_train, y_test = train_test_split(normal_counts.iloc[0:,1:].
      ↪T, meta_data['Simplified_class'], test_size=0.3, random_state = 10101)
```

```
[10]: X_train.head()
```

```
[10]:
```

	0	1	2	3	4	5	\
DLDR_0036	5.820135	-1.060061	4.388400	4.080172	2.564430	3.552685	
DLDR_0081	6.546299	0.582165	3.752090	4.645175	3.840899	3.201075	
DLDR_0191	6.604050	-0.865036	4.351489	4.072137	3.143138	4.037476	
DLDR_0188	6.480745	-1.083676	4.361634	4.313540	3.120196	1.941859	
DLDR_0130	6.550016	-1.222374	4.534941	4.370763	3.512952	2.517867	

	6	7	8	9	...	17386	17387	\
DLDR_0036	11.011379	4.682305	6.951539	4.589555	...	4.851631	-0.719024	
DLDR_0081	11.433579	3.705547	7.143316	5.482169	...	5.050934	1.681701	
DLDR_0191	11.782524	4.358527	8.468526	5.013154	...	5.112244	-2.002540	
DLDR_0188	11.451981	4.556052	7.183779	5.066071	...	5.204366	2.086249	
DLDR_0130	12.041229	4.315374	7.485927	4.655817	...	5.519093	-0.037949	

	17388	17389	17390	17391	17392	17393	\
DLDR_0036	-0.719024	-1.266512	-0.323095	3.157170	0.459313	3.115198	
DLDR_0081	-0.640228	0.096738	0.582165	4.041596	0.096738	2.418666	
DLDR_0191	-0.076540	-1.154543	-0.624028	2.050571	-1.517113	2.084923	
DLDR_0188	0.501286	0.632531	-0.431600	2.616763	1.063165	1.941859	
DLDR_0130	-0.289488	-1.874450	-0.981365	2.594319	0.447478	2.556600	

	17394	17395
DLDR_0036	-2.645023	0.760969
DLDR_0081	-2.225190	0.582165
DLDR_0191	-4.324468	0.067849
DLDR_0188	-2.306069	0.632531
DLDR_0130	-2.359877	0.932905

[5 rows x 17396 columns]

```
[11]: X_train.shape
```

```
[11]: (134, 17396)
```

```
[12]: y_train.head()
```

```
[12]: 35          Normal
      80    Advanced_fibrosis
      190         Normal
      187         Normal
      129  Non_advanced_Fibrosis
      Name: Simplified_class, dtype: object
```

```
[13]: y_train.shape
```

```
[13]: (134,)
```

Labels distribution in train set:

```
[14]: print("Number of Normal samples: ", y_train.to_list().count("Normal"))
      print("Number of Non_advanced_Fibrosis samples: ", y_train.to_list().
      ↪count("Non_advanced_Fibrosis"))
      print("Number of Advanced_fibrosis samples: ", y_train.to_list().
      ↪count("Advanced_fibrosis"))
```

```
Number of Normal samples: 44
Number of Non_advanced_Fibrosis samples: 38
Number of Advanced_fibrosis samples: 52
```

Now, let's save the training set data to use in the R Jupyter Notebook for feature selection

```
[15]: X_train.T.to_csv('train_normal_counts.csv', index=False)
      y_train.T.to_csv('train_meta_data.csv', index=False)
```

Run R Jupyter Notebook Let's load the output of R to continue the task

```
[16]: subset_data = pd.read_csv('subset_data.csv')
```

```
[17]: subset_data
```

```
[17]: Unnamed: 0  DLDR_0036  DLDR_0081  DLDR_0191  DLDR_0188  DLDR_0130  \
0           10    4.589555    5.482169    5.013154    5.066071    4.655817
1           57   -0.719024    0.582165   -1.154543   -1.083676   -0.289488
2          265    0.587637   -2.225190   -0.076540    0.196432    0.272391
3          275    1.662405    2.529697    1.508422    1.991612    1.467942
4          278    5.529902    5.846272    5.553583    5.359267    5.774549
..          ...          ...          ...          ...          ...          ...
```

522	16863	3.811126	3.329399	2.183327	2.412750	3.323819	
523	16887	3.576080	2.418666	3.355012	3.288878	3.231084	
524	16892	4.007463	3.060212	2.951656	3.366357	2.649112	
525	17075	6.925971	7.323632	7.411511	7.189120	7.519400	
526	17187	-2.645023	-2.225190	-4.324468	-3.891031	-4.681805	
	DLDR_0013	DLDR_0079	DLDR_0131	DLDR_0135	...	DLDR_0175	DLDR_0052 \
0	4.299078	4.752957	5.409514	4.993777	...	4.853565	4.567173
1	-0.006469	-1.138550	-4.060128	-2.104255	...	-2.285986	-0.110595
2	2.523150	1.424386	0.332190	0.217673	...	1.083248	0.352377
3	1.934152	2.031375	2.048397	0.528013	...	1.083248	2.556469
4	4.962662	5.201300	5.500205	5.782051	...	5.862067	5.038603
..
522	2.690386	3.515760	3.197260	2.112975	...	1.555316	3.901539
523	3.757011	2.435441	3.079423	2.393996	...	2.974541	3.737944
524	2.774044	3.123117	2.951099	2.974696	...	2.604785	3.934981
525	6.661844	7.087205	7.366661	7.586092	...	7.491543	6.721953
526	0.820694	-3.013019	-4.060128	-4.426183	...	-3.870949	-3.735085
	DLDR_0087	DLDR_0155	DLDR_0092	DLDR_0187	DLDR_0186	DLDR_0179	\
0	4.985158	5.234776	4.900754	4.923087	5.004500	5.058062	
1	-0.385647	-3.580073	0.317259	-0.757399	-3.486017	-1.346471	
2	-0.309698	0.120366	0.890444	-0.291735	0.811664	-0.246936	
3	2.055951	1.374123	1.902221	1.838969	1.643266	1.272439	
4	5.665231	5.781140	5.745159	5.438062	5.570169	5.558701	
..
522	3.344479	1.911780	3.763838	3.667268	2.742802	2.446087	
523	2.476362	2.817958	2.596331	3.203431	3.162640	2.353968	
524	3.001089	2.596515	3.945549	4.002569	3.359473	3.497226	
525	7.511299	7.684565	7.620817	7.419428	7.417614	7.354464	
526	-4.010138	-5.165036	-2.718365	-5.149716	-5.070979	-2.568864	
	DLDR_0182	DLDR_0001					
0	5.002454	4.221450					
1	-1.103280	-0.005377					
2	0.416094	2.729099					
3	1.932344	2.467558					
4	5.733822	5.220121					
..					
522	3.857108	2.844001					
523	3.445156	3.271068					
524	3.563476	2.830128					
525	7.633460	6.752731					
526	-3.425208	1.202731					

[527 rows x 135 columns]

```
[18]: subset_data.shape
```

```
[18]: (527, 135)
```

We need to extract the same features for the X test dataset as well. To do this, we must identify the indices of our selected DEGs and apply the same subsetting to the test data.

```
[19]: selected_genes_R = subset_data.T.iloc[0,:].to_list()
```

```
[20]: selected_genes_Python = [int(i-1) for i in selected_genes_R ]
```

```
[21]: selected_genes_Python[:10]
```

```
[21]: [9, 56, 264, 274, 277, 296, 309, 340, 351, 389]
```

```
[22]: len(selected_genes_Python)
```

```
[22]: 527
```

```
[23]: df_deg = subset_data.iloc[0:,1:].T  
df_deg_test = X_test[selected_genes_Python]
```

```
[24]: df_deg.head()
```

```
[24]:
```

	0	1	2	3	4	5	\
DLDR_0036	4.589555	-0.719024	0.587637	1.662405	5.529902	7.549487	
DLDR_0081	5.482169	0.582165	-2.225190	2.529697	5.846272	7.759228	
DLDR_0191	5.013154	-1.154543	-0.076540	1.508422	5.553583	7.612538	
DLDR_0188	5.066071	-1.083676	0.196432	1.991612	5.359267	8.070057	
DLDR_0130	4.655817	-0.289488	0.272391	1.467942	5.774549	7.439405	

	6	7	8	9	...	517	518	\
DLDR_0036	0.318451	6.089009	5.898008	9.044188	...	3.104511	1.602904	
DLDR_0081	0.096738	6.539682	6.101239	8.633568	...	2.819204	1.234242	
DLDR_0191	1.033084	6.434588	6.098648	8.244201	...	2.275445	0.884985	
DLDR_0188	0.356896	6.476384	6.118797	7.951712	...	2.708882	0.356896	
DLDR_0130	0.362589	6.486240	6.239292	8.248378	...	2.969247	1.248932	

	519	520	521	522	523	524	\
DLDR_0036	0.915692	7.886943	5.319606	3.811126	3.576080	4.007463	
DLDR_0081	-0.640228	6.310085	3.797178	3.329399	2.418666	3.060212	
DLDR_0191	0.067849	7.139567	3.909152	2.183327	3.355012	2.951656	
DLDR_0188	-0.431600	7.167637	4.021858	2.412750	3.288878	3.366357	
DLDR_0130	-1.222374	6.106913	3.570860	3.323819	3.231084	2.649112	

	525	526
DLDR_0036	6.925971	-2.645023
DLDR_0081	7.323632	-2.225190

```
DLDR_0191  7.411511 -4.324468
DLDR_0188  7.189120 -3.891031
DLDR_0130  7.519400 -4.681805
```

[5 rows x 527 columns]

```
[25]: df_deg_test.head()
```

```
[25]:
```

	9	56	264	274	277	296	\
DLDR_0022	3.967031	0.177378	2.880522	2.254454	5.101756	7.696151	
DLDR_0016	4.603126	0.002788	2.608214	1.992446	5.079298	7.453528	
DLDR_0004	4.003661	-0.178864	2.418699	2.287454	5.125878	6.943351	
DLDR_0165	5.286333	0.143691	0.332137	1.587791	5.827910	7.489948	
DLDR_0127	5.118729	-0.818750	0.897457	1.530400	5.446251	7.622416	

	309	340	351	389	...	16641	16686	\
DLDR_0022	2.077214	5.845657	5.444621	8.076772	...	2.825450	0.246091	
DLDR_0016	2.274384	5.735367	5.394124	9.027927	...	2.934201	0.299770	
DLDR_0004	2.538993	6.049955	5.583037	8.243201	...	2.418699	0.217065	
DLDR_0165	1.225221	6.994423	6.261908	8.733452	...	2.368478	0.143691	
DLDR_0127	0.897457	6.472366	5.832644	8.246698	...	2.137661	0.811301	

	16715	16731	16762	16862	16886	16891	\
DLDR_0022	0.928350	7.735892	4.644785	2.335526	3.460071	3.723708	
DLDR_0016	0.488215	7.956485	5.527246	2.274384	3.722128	3.226652	
DLDR_0004	-0.519901	7.122512	3.978350	2.974088	3.280568	2.418699	
DLDR_0165	0.389852	6.695064	3.291074	1.485611	2.862085	2.892352	
DLDR_0127	1.128783	6.668899	3.631283	3.308569	2.881690	2.101816	

	17074	17186
DLDR_0022	6.646771	0.246091
DLDR_0016	6.739625	0.545930
DLDR_0004	6.827647	0.999473
DLDR_0165	7.434889	-3.697611
DLDR_0127	7.391864	0.215197

[5 rows x 527 columns]

It's time for classification. I have used four types of classifiers throughout the project to ensure consistency. These classifiers are well-established in machine learning and are well-suited for this task. The classifiers used are: Logistic Regression, Support Vector Machine, Random Forest, and Multilayer Perceptron.

Logistic Regression

```
[26]: # Logistic Regression
LR_model = LogisticRegression(solver='saga')
LR_model.fit(df_deg, y_train)
```

```

y_pred_LR = LR_model.predict(df_deg_test)
LR_precision = precision_score(y_test, y_pred_LR, average='macro')
LR_recall = recall_score(y_test, y_pred_LR, average='macro')
print(LR_precision)
print(LR_recall)

```

0.8010912698412698
0.7752136752136751

Support Vector Machine

```

[27]: # Support Vector Machine
SVM_model = SVC(kernel='linear', C=1)
SVM_model.fit(df_deg, y_train)
y_pred_SVM = SVM_model.predict(df_deg_test)
SVM_precision = precision_score(y_test, y_pred_SVM, average='macro')
SVM_recall = recall_score(y_test, y_pred_SVM, average='macro')
print(SVM_precision)
print(SVM_recall)

```

0.8441558441558442
0.8008547008547008

Random Forest

```

[28]: # Random Forest
RF_model = RandomForestClassifier(random_state = 10101)
RF_model.fit(df_deg, y_train)
y_pred_RF = RF_model.predict(df_deg_test)
RF_precision = precision_score(y_test, y_pred_RF, average='macro')
RF_recall = recall_score(y_test, y_pred_RF, average='macro')
print(RF_precision)
print(RF_recall)

```

0.7743589743589744
0.7606837606837606

Multi Layer Perceptron

```

[29]: # Multi Layer Perceptron
MLP_model = MLPClassifier(hidden_layer_sizes=(100,), max_iter=300,
    ↪activation='relu', solver='adam', random_state = 10101)
MLP_model.fit(df_deg, y_train)
y_pred_MLP = MLP_model.predict(df_deg_test)
MLP_precision = precision_score(y_test, y_pred_MLP, average='macro')
MLP_recall = recall_score(y_test, y_pred_MLP, average='macro')
print(MLP_precision)
print(MLP_recall)

```

0.7317550505050505
0.7051282051282052

So far, we have obtained only a single result for each classifier, but these results are not statistically valid. To test for reproducibility, we have put the entire process into a loop to run it 100 times. Additionally, we created an R script to perform DE analysis automatically using subprocess. Let's run this cell and save the results for statistical analysis.

```
[30]: import subprocess

# Load Data
normal_counts = pd.read_csv('Normal_counts.voom.csv')
meta_data = pd.read_csv('meta_data.csv')

n_iterations = 300
test_size = 0.3

LR_precisions = []
LR_recalls = []

SVM_precisions = []
SVM_recalls = []

RF_precisions = []
RF_recalls = []

MLP_precisions = []
MLP_recalls = []

for i in range(100):
    print('iteration',i)
    X_train, X_test, y_train, y_test = train_test_split(normal_counts.iloc[0:,1:
↪].T, meta_data['Simplified_class'], test_size=0.3, random_state = i)
    X_train.T.to_csv('train_normal_counts.csv', index=False)
    y_train.T.to_csv('train_meta_data.csv', index=False)

    r_script_path = r"q1r.R"
    rsctest_path = r"C:\Program Files\R\R-4.2.1\bin\Rscript.exe"
    # Execute the R script
    try:
        subprocess.run([rsctest_path, r_script_path], capture_output=True,
↪text=True)
    except subprocess.CalledProcessError as e:
        print(f"Error executing R script: {e}")

    #
    subset_data = pd.read_csv('subset_data.csv')
    selected_genes_R = subset_data.T.iloc[0,:].to_list()
    selected_genes_Python = [int(i-1) for i in selected_genes_R ]
    df_deg = subset_data.iloc[0:,1:].T
```

```

df_deg_test = X_test[selected_genes_Python]

# Logistic Regression
LR_model = LogisticRegression(solver='saga')
LR_model.fit(df_deg, y_train)
y_pred_LR = LR_model.predict(df_deg_test)
LR_precision = precision_score(y_test, y_pred_LR, average='macro')
LR_recall = recall_score(y_test, y_pred_LR, average='macro')
LR_precisions.append(LR_precision)
LR_recalls.append(LR_recall)

# Support Vector Machine
SVM_model = SVC(kernel='linear', C=1)
SVM_model.fit(df_deg, y_train)
y_pred_SVM = SVM_model.predict(df_deg_test)
SVM_precision = precision_score(y_test, y_pred_SVM, average='macro')
SVM_recall = recall_score(y_test, y_pred_SVM, average='macro')
SVM_precisions.append(SVM_precision)
SVM_recalls.append(SVM_recall)

# Random Forest
RF_model = RandomForestClassifier(random_state = i)
RF_model.fit(df_deg, y_train)
y_pred_RF = RF_model.predict(df_deg_test)
RF_precision = precision_score(y_test, y_pred_RF, average='macro')
RF_recall = recall_score(y_test, y_pred_RF, average='macro')
RF_precisions.append(RF_precision)
RF_recalls.append(RF_recall)

# Multi Layer Perceptron
MLP_model = MLPClassifier(hidden_layer_sizes=(100,), max_iter=300,
activation='relu', solver='adam', random_state = i)
MLP_model.fit(df_deg, y_train)
y_pred_MLP = MLP_model.predict(df_deg_test)
MLP_precision = precision_score(y_test, y_pred_MLP, average='macro')
MLP_recall = recall_score(y_test, y_pred_MLP, average='macro')
MLP_precisions.append(MLP_precision)
MLP_recalls.append(MLP_recall)

```

```

iteration 0
iteration 1
iteration 2
iteration 3
iteration 4
iteration 5
iteration 6
iteration 7

```

iteration 8
iteration 9
iteration 10
iteration 11
iteration 12
iteration 13
iteration 14
iteration 15
iteration 16
iteration 17
iteration 18
iteration 19
iteration 20
iteration 21
iteration 22
iteration 23
iteration 24
iteration 25
iteration 26
iteration 27
iteration 28
iteration 29
iteration 30
iteration 31
iteration 32
iteration 33
iteration 34
iteration 35
iteration 36
iteration 37
iteration 38
iteration 39
iteration 40
iteration 41
iteration 42
iteration 43
iteration 44
iteration 45
iteration 46
iteration 47
iteration 48
iteration 49
iteration 50
iteration 51
iteration 52
iteration 53
iteration 54
iteration 55

iteration 56
iteration 57
iteration 58
iteration 59
iteration 60
iteration 61
iteration 62
iteration 63
iteration 64
iteration 65
iteration 66
iteration 67
iteration 68
iteration 69
iteration 70
iteration 71
iteration 72
iteration 73
iteration 74
iteration 75
iteration 76
iteration 77
iteration 78
iteration 79
iteration 80
iteration 81
iteration 82
iteration 83
iteration 84
iteration 85
iteration 86
iteration 87
iteration 88
iteration 89
iteration 90
iteration 91
iteration 92
iteration 93
iteration 94
iteration 95
iteration 96
iteration 97
iteration 98
iteration 99

Last but not least, we need to report the average and confidence interval for the results of each classifier. We will also create a bar plot to visually compare the results of the classifiers.

```
[31]: LR_mean_precision = np.mean(LR_precisions)
LR_mean_recall = np.mean(LR_recalls)
LR_precision_conf_interval = np.percentile(LR_precisions, [2.5, 97.5])
LR_recall_conf_interval = np.percentile(LR_recalls, [2.5, 97.5])
print(f'Mean Precision for Logistic Regression: {LR_mean_precision}, 95% CI:␣
↪{LR_precision_conf_interval}')
print(f'Mean Recall for Logistic Regression: {LR_mean_recall}, 95% CI:␣
↪{LR_recall_conf_interval}')
print("-----")

SVM_mean_precision = np.mean(SVM_precisions)
SVM_mean_recall = np.mean(SVM_recalls)
SVM_precision_conf_interval = np.percentile(SVM_precisions, [2.5, 97.5])
SVM_recall_conf_interval = np.percentile(SVM_recalls, [2.5, 97.5])
print(f'Mean Precision for Support Vector Machine: {SVM_mean_precision}, 95% CI:
↪ {SVM_precision_conf_interval}')
print(f'Mean Recall for Support Vector Machine: {SVM_mean_recall}, 95% CI:␣
↪{SVM_recall_conf_interval}')
print("-----")

RF_mean_precision = np.mean(RF_precisions)
RF_mean_recall = np.mean(RF_recalls)
RF_precision_conf_interval = np.percentile(RF_precisions, [2.5, 97.5])
RF_recall_conf_interval = np.percentile(RF_recalls, [2.5, 97.5])
print(f'Mean Precision for Random Forest: {RF_mean_precision}, 95% CI:␣
↪{RF_precision_conf_interval}')
print(f'Mean Recall for Random Forest: {RF_mean_recall}, 95% CI:␣
↪{RF_recall_conf_interval}')
print("-----")

MLP_mean_precision = np.mean(MLP_precisions)
MLP_mean_recall = np.mean(MLP_recalls)
MLP_precision_conf_interval = np.percentile(MLP_precisions, [2.5, 97.5])
MLP_recall_conf_interval = np.percentile(MLP_recalls, [2.5, 97.5])
print(f'Mean Precision for Multi Layer Perceptron: {MLP_mean_precision}, 95% CI:
↪ {MLP_precision_conf_interval}')
print(f'Mean Recall for Multi Layer Perceptron: {MLP_mean_recall}, 95% CI:␣
↪{MLP_recall_conf_interval}')
```

Mean Precision for Logistic Regression: 0.7636803304888079, 95% CI: [0.64076229
0.85890152]

Mean Recall for Logistic Regression: 0.7645416569253737, 95% CI: [0.63156272
0.87401786]

Mean Precision for Support Vector Machine: 0.7743040927880918, 95% CI:
[0.66116453 0.85674319]

Mean Recall for Support Vector Machine: 0.7723320280185707, 95% CI: [0.65753472 0.86513971]

Mean Precision for Random Forest: 0.7860810773240244, 95% CI: [0.6725455 0.87525208]

Mean Recall for Random Forest: 0.7867412414621653, 95% CI: [0.65607372 0.88233974]

Mean Precision for Multi Layer Perceptron: 0.7677514303520464, 95% CI: [0.65656233 0.85438582]

Mean Recall for Multi Layer Perceptron: 0.7618658834643532, 95% CI: [0.6144042 0.8572071]

```
[32]: model_results = {
    'RandomForest': {
        'mean_precision': RF_mean_precision,
        'precision_ci': RF_precision_conf_interval,
        'mean_recall': RF_mean_recall,
        'recall_ci': RF_precision_conf_interval
    },
    'SVM': {
        'mean_precision': SVM_mean_precision,
        'precision_ci': SVM_precision_conf_interval,
        'mean_recall': SVM_mean_recall,
        'recall_ci': SVM_precision_conf_interval
    },
    'LogisticRegression': {
        'mean_precision': LR_mean_precision,
        'precision_ci': LR_precision_conf_interval,
        'mean_recall': LR_mean_recall,
        'recall_ci': LR_precision_conf_interval
    },
    'MLP': {
        'mean_precision': MLP_mean_precision,
        'precision_ci': MLP_precision_conf_interval,
        'mean_recall': MLP_mean_recall,
        'recall_ci': MLP_precision_conf_interval
    }
}
```

```
[33]: models = list(model_results.keys())
mean_precisions = [model_results[model]['mean_precision'] for model in models]
precision_cis = [model_results[model]['precision_ci'] for model in models]
mean_recalls = [model_results[model]['mean_recall'] for model in models]
recall_cis = [model_results[model]['recall_ci'] for model in models]
```


```
precision_errors = np.array([[mean - ci[0], ci[1] - mean] for mean, ci in
    ↪zip(mean_precisions, precision_cis)]).T
recall_errors = np.array([[mean - ci[0], ci[1] - mean] for mean, ci in
    ↪zip(mean_recalls, recall_cis)]).T
```

```
[34]: fig, ax = plt.subplots(1, 2, figsize=(14, 6))

# Plot precision
ax[0].bar(models, mean_precisions, yerr=precision_errors, capsize=5,
    ↪color='skyblue')
ax[0].set_title('Mean Precision with 95% Confidence Interval')
ax[0].set_ylabel('Precision')
ax[0].set_ylim([0.5, 1])
ax[0].set_xticklabels(models, rotation=45, ha="right")

# Plot recall
ax[1].bar(models, mean_recalls, yerr=recall_errors, capsize=5,
    ↪color='lightgreen')
ax[1].set_title('Mean Recall with 95% Confidence Interval')
ax[1].set_ylabel('Recall')
ax[1].set_ylim([0.5, 1])
ax[1].set_xticklabels(models, rotation=45, ha="right")

plt.tight_layout()
plt.show()
```



Here are the final results: as you can see, the performance of the classifiers is very similar, with none of them significantly outperforming the others. However, Random Forest exhibits slightly higher precision and recall compared to the other models.