

Q4-Python

August 15, 2024

1 NAFLD detection using RNA-Seq data

Content creators: Mahdi Anvari, Sadegh Rizi

University of Tehran, Department of Biotechnology

In this part, we need to perform a regression task to predict the age of samples based on their gene expression profiles and clinical data.

We will then split the dataset based on sex, repeat the regression task, and compare the results between the groups.

2 Q4 - Python section

```
[24]: # Import needed libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.metrics import precision_score, recall_score
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder
from sklearn.linear_model import LassoCV
import warnings
warnings.filterwarnings("ignore")
from sklearn.linear_model import LinearRegression, Ridge, ElasticNet
from sklearn.svm import SVR
from sklearn.metrics import mean_squared_error, r2_score
from sklearn.preprocessing import StandardScaler, OneHotEncoder
from sklearn.compose import ColumnTransformer
from sklearn.pipeline import Pipeline
from sklearn.feature_selection import VarianceThreshold
```

```
[2]: # Load Data
normal_counts = pd.read_csv('Normal.counts.voom.csv')
meta_data = pd.read_csv('meta_data.csv')
```

```
[3]: normal_counts.head()
```

```
[3]:
```

	gene	DLDR_0001	DLDR_0002	DLDR_0003	DLDR_0004	DLDR_0005	\
0	ENSG000000000003	5.965571	5.741587	5.996891	5.551919	6.430237	
1	ENSG000000000005	1.612375	2.147793	0.418542	0.702492	1.215978	
2	ENSG000000000419	4.133821	4.120969	4.086129	4.116240	4.393797	
3	ENSG000000000457	4.111056	3.922234	3.964871	3.978350	4.018235	
4	ENSG000000000460	4.150662	3.732756	3.634637	3.853979	3.614220	

	DLDR_0006	DLDR_0007	DLDR_0008	DLDR_0009	...	DLDR_0183	DLDR_0184	\
0	6.234619	6.071503	6.441882	5.752712	...	6.304802	6.576246	
1	0.920810	0.458163	0.927224	1.089389	...	-0.031596	-1.091275	
2	4.390909	4.148242	4.554655	4.203819	...	4.176599	4.244459	
3	3.864521	4.263119	3.896271	4.139546	...	4.378342	4.453838	
4	3.500857	4.000565	4.016287	3.904500	...	2.974209	3.720038	

	DLDR_0185	DLDR_0186	DLDR_0187	DLDR_0188	DLDR_0189	DLDR_0190	\
0	6.735760	6.344234	6.608924	6.480745	6.360397	6.367705	
1	-0.942637	-0.026585	-0.757399	-1.083676	0.886550	-0.902201	
2	4.342765	4.179319	4.274450	4.361634	4.093280	4.148010	
3	4.685598	4.438796	4.042577	4.313540	4.205119	4.506058	
4	4.640011	3.814717	2.126408	3.120196	3.336802	3.982071	

	DLDR_0191	DLDR_0192
0	6.604050	6.514539
1	-0.865036	-1.588749
2	4.351489	3.859711
3	4.072137	4.341988
4	3.143138	2.741172

[5 rows x 193 columns]

```
[4]: normal_counts.shape
```

```
[4]: (17396, 193)
```

```
[5]: meta_data.head()
```

```
[5]:
```

	Patient_ID	SEX	BMI_surg	Age	Run	Diabet	\
0	DLDR_0001	Female	35.214555	55	SRR8378590	Non Diabetic	
1	DLDR_0002	Female	39.421748	47	SRR8378589	Diabetic	
2	DLDR_0003	Male	48.758108	46	SRR8378432	Non Diabetic	
3	DLDR_0004	Female	41.822607	36	SRR8378431	Non Diabetic	
4	DLDR_0005	Female	53.582192	54	SRR8378434	Non Diabetic	

	Simplified_class
0	Normal
1	Normal
2	Normal

```
3         Normal
4         Normal
```

```
[6]: meta_data.shape
```

```
[6]: (192, 7)
```

```
[13]: print("Number of Female samples: ", meta_data['SEX'].to_list().count("Female"))
      print("Number of Male samples: ", meta_data['SEX'].to_list().count("Male"))
```

```
Number of Female samples: 154
```

```
Number of Male samples: 38
```

To include our clinical data, we need to perform some preprocessing steps. For categorical data, we will use one-hot encoding, and for continuous data, we will apply normalization.

```
[14]: # One-Hot Encoding for categorical features
      categorical_features = ['SEX', 'Diabet']
      ohe = OneHotEncoder()
      # Standardize continuous variables
      numerical_features = ['BMI_surg', 'Age']
      scaler = StandardScaler()
      # Preprocess clinical data
      preprocessor = ColumnTransformer(
          transformers=[
              ('num', scaler, numerical_features),
              ('cat', ohe, categorical_features)])
```

```
[15]: clinical_data = preprocessor.fit_transform(meta_data)
```

```
[16]: clinical_data = pd.DataFrame(clinical_data).iloc[:, [0, 2, 4]]
      clinical_data.columns = ['BMI_surg', 'Sex', 'Diabet']
```

```
[17]: clinical_data
```

```
[17]:   BMI_surg  Sex  Diabet
0  -1.308155  1.0    0.0
1  -0.857325  1.0    1.0
2   0.143131  0.0    0.0
3  -0.600056  1.0    0.0
4   0.660065  1.0    0.0
..      ...  ...    ...
187 -0.772196  0.0    1.0
188  0.284524  1.0    0.0
189  6.161411  0.0    1.0
190  0.007788  1.0    0.0
191 -0.214761  1.0    0.0
```

[192 rows x 3 columns]

```
[18]: gene_names = normal_counts.iloc[:,0]
gene_names.at[17396] = 'BMI_surg'
gene_names.at[17397] = 'SEX'
gene_names.at[17398] = 'Diabet'
```

```
[19]: gene_names
```

```
[19]: 0      ENSG000000000003
1      ENSG000000000005
2      ENSG000000000419
3      ENSG000000000457
4      ENSG000000000460
...
17394  ENSG00000273274
17395  ENSG00000273294
17396      BMI_surg
17397      SEX
17398      Diabet
Name: gene, Length: 17399, dtype: object
```

Now, let's combine our gene expression data with clinical data

```
[20]: combined_data = np.hstack((normal_counts.iloc[0:,1:].T, clinical_data))
df_deg = pd.DataFrame(combined_data)
df_deg.columns = gene_names
```

```
[21]: df_deg
```

```
[21]: gene  ENSG000000000003  ENSG000000000005  ENSG000000000419  ENSG000000000457  \
0      5.965571      1.612375      4.133821      4.111056
1      5.741587      2.147793      4.120969      3.922234
2      5.996891      0.418542      4.086129      3.964871
3      5.551919      0.702492      4.116240      3.978350
4      6.430237      1.215978      4.393797      4.018235
..      ...
187     6.480745     -1.083676      4.361634      4.313540
188     6.360397      0.886550      4.093280      4.205119
189     6.367705     -0.902201      4.148010      4.506058
190     6.604050     -0.865036      4.351489      4.072137
191     6.514539     -1.588749      3.859711      4.341988

gene  ENSG000000000460  ENSG000000000938  ENSG000000000971  ENSG00000001036  \
0      4.150662      2.975845      11.005488      4.405768
1      3.732756      3.199989      10.860700      3.895350
2      3.634637      2.949733      10.934025      4.282577
3      3.853979      2.991061      10.760445      4.297722
```

4	3.614220	2.836130	11.491427	4.405558
..
187	3.120196	1.941859	11.451981	4.556052
188	3.336802	2.911496	11.658711	4.315153
189	3.982071	2.202136	11.864397	4.348255
190	3.143138	4.037476	11.782524	4.358527
191	2.741172	2.803568	11.474777	3.895962

gene	ENSG00000001084	ENSG00000001167	...	ENSG00000272658	ENSG00000272869	\
0	6.825329	4.221450	...	1.365460	0.691230	
1	6.453687	4.218183	...	0.935490	1.004202	
2	6.437658	3.736947	...	0.675699	0.943633	
3	6.710840	4.003661	...	0.702492	0.431190	
4	7.437655	4.377965	...	0.898496	0.085581	
..	
187	7.183779	5.066071	...	0.632531	-0.431600	
188	7.384726	4.627703	...	-0.147398	-1.080284	
189	7.326297	4.528530	...	0.048889	-1.108652	
190	8.468526	5.013154	...	-1.154543	-0.624028	
191	7.617180	4.778544	...	-0.114818	-0.510747	

gene	ENSG00000273079	ENSG00000273173	ENSG00000273259	ENSG00000273274	\
0	5.637483	-0.005377	1.675800	2.683536	
1	5.975612	0.532134	1.555218	2.926666	
2	5.531648	-0.184123	2.391906	2.260662	
3	5.571799	-0.034474	1.639298	2.341393	
4	5.636848	-1.216981	1.974160	1.351861	
..	
187	2.616763	1.063165	1.941859	-2.306069	
188	2.975985	0.589568	2.754906	-0.839275	
189	3.218814	-0.416774	2.246191	-1.639167	
190	2.050571	-1.517113	2.084923	-4.324468	
191	2.572243	0.667591	2.160189	-3.173712	

gene	ENSG00000273294	BMI_surg	SEX	Diabet
0	-0.339797	-1.308155	1.0	0.0
1	0.435919	-0.857325	1.0	1.0
2	-0.691083	0.143131	0.0	0.0
3	0.096771	-0.600056	1.0	0.0
4	-0.079478	0.660065	1.0	0.0
..
187	0.632531	-0.772196	0.0	1.0
188	-0.016153	0.284524	1.0	0.0
189	0.145105	6.161411	0.0	1.0
190	0.067849	0.007788	1.0	0.0
191	1.218606	-0.214761	1.0	0.0

[192 rows x 17399 columns]

Let's split our data and labels to train and test sets

```
[22]: X_train, X_test, y_train, y_test = train_test_split(df_deg, meta_data['Age'],  
↳ test_size=0.2, random_state = 10101)
```

First, we applied a variance threshold to filter out features that do not have sufficient variation.

```
[25]: threshold = 0.8  
selector = VarianceThreshold(threshold)  
filtered_data = selector.fit_transform(X_train)  
selected_features = X_train.columns[selector.get_support()]  
X_train_filtered = pd.DataFrame(filtered_data, columns=selected_features)
```

```
[26]: X_train_filtered.shape
```

```
[26]: (153, 5229)
```

Next, we performed LASSO to extract the most relevant features for this task.

```
[21]: lasso = LassoCV(cv=5).fit(X_train_filtered, y_train)  
# Extract the non-zero coefficients  
important_genes = X_train_filtered.columns[(lasso.coef_ != 0)]
```

```
[22]: important_genes
```

```
[22]: Index(['ENSG000000001630', 'ENSG000000007001', 'ENSG000000011083',  
          'ENSG000000024526', 'ENSG000000064205', 'ENSG000000071909',  
          'ENSG000000086548', 'ENSG000000091129', 'ENSG000000103154',  
          'ENSG000000104369', 'ENSG000000109182', 'ENSG000000115009',  
          'ENSG000000115457', 'ENSG000000129824', 'ENSG000000131080',  
          'ENSG000000134184', 'ENSG000000141668', 'ENSG000000144229',  
          'ENSG000000151025', 'ENSG000000152726', 'ENSG000000152779',  
          'ENSG000000162391', 'ENSG000000164266', 'ENSG000000165246',  
          'ENSG000000166923', 'ENSG000000170160', 'ENSG000000171496',  
          'ENSG000000171747', 'ENSG000000181333', 'ENSG000000182771',  
          'ENSG000000182870', 'ENSG000000184674', 'ENSG000000184698',  
          'ENSG000000189266', 'ENSG000000221890', 'ENSG000000236882',  
          'ENSG000000244694', 'ENSG000000255798', 'ENSG000000256618'],  
          dtype='object', name='gene')
```

```
[23]: len(important_genes)
```

```
[23]: 39
```

```
[24]: X_train[important_genes]
```

```

[24]: gene ENSG00000001630 ENSG00000007001 ENSG00000011083 ENSG00000024526 \
9 -0.579836 5.686951 2.307689 0.467470
16 0.467219 3.764056 3.043866 2.238594
149 -1.395729 6.789909 -5.096169 -1.008706
63 -0.910103 2.403314 2.927247 1.445378
41 0.522448 6.611324 -4.969405 -4.969405
.. ...
186 -0.105322 7.730441 -5.149716 -5.149716
185 -0.823052 7.831961 -3.486017 -1.370540
178 -0.983901 7.563422 -2.568864 -2.568864
181 -0.255283 4.830607 -5.010171 -5.010171
0 -0.956468 3.061454 2.537521 1.019285

gene ENSG000000064205 ENSG000000071909 ENSG000000086548 ENSG000000091129 \
9 1.067862 1.780060 0.036835 3.584736
16 2.238594 3.534050 2.575022 3.603591
149 -2.288814 -0.703852 -0.238188 0.685191
63 1.770618 2.803334 1.991119 2.777221
41 -1.268966 -4.969405 -1.509974 0.456859
.. ...
186 -3.564753 -0.901788 -5.149716 -1.242825
185 -5.070979 -3.486017 -2.749051 0.601446
178 -1.831898 -4.153826 -4.153826 -0.246936
181 -2.688243 -2.688243 -2.688243 -1.550739
0 1.823142 2.924544 2.234674 2.937541

gene ENSG00000103154 ENSG00000104369 ... ENSG00000182771 ENSG00000182870 \
9 3.075516 -1.027295 ... 2.052432 2.254740
16 3.027090 1.223948 ... 2.957966 2.575022
149 2.985980 -1.636737 ... 2.355042 -5.096169
63 3.635429 0.933171 ... 2.533816 2.549329
41 3.235166 -4.969405 ... 1.096684 -3.384443
.. ...
186 1.536785 -2.827788 ... -2.827788 -3.564753
185 1.774511 -1.901054 ... 2.047962 -5.070979
178 1.868542 -1.831898 ... 0.800370 -1.831898
181 2.407682 -2.202816 ... -1.550739 -3.425208
0 2.911429 0.691230 ... 2.975845 2.055505

gene ENSG00000184674 ENSG00000184698 ENSG00000189266 ENSG00000221890 \
9 4.527294 -0.399264 0.467470 0.722727
16 3.659086 0.894640 -0.142835 2.267449
149 3.889673 -2.288814 -2.774241 1.503744
63 5.584576 0.674859 -4.997566 1.822613
41 -4.969405 -2.647477 1.659951 -2.647477
.. ...
186 4.975697 -5.149716 0.827564 0.059737

```

185	-5.070979	-1.611548	2.005836	-0.316092
178	3.782812	-2.568864	0.800370	-0.246936
181	3.644465	-5.010171	1.619186	-1.840246
0	5.048033	0.751351	1.644436	1.706497

gene	ENSG000000236882	ENSG000000244694	ENSG000000255798	ENSG000000256618
9	6.884506	0.798676	2.432137	6.505504
16	5.158060	2.209151	-0.690322	7.458252
149	4.096124	-1.395729	-5.096169	4.302575
63	5.936862	1.411825	1.541593	6.323799
41	8.626668	-2.647477	-2.162050	6.275553
..
186	7.257286	-1.242825	-1.062253	9.595696
185	4.631193	-2.263624	2.148189	8.637566
178	7.818077	-1.831898	-2.568864	7.130998
181	4.988419	0.481682	-1.840246	9.887060
0	3.960009	1.245166	2.007006	6.124616

[153 rows x 39 columns]

```
[25]: X_test[important_genes].shape
```

```
[25]: (39, 39)
```

It's time for regression. I have used three types of regressors in this part to ensure consistency. These regressors are well-established in machine learning and are well-suited for this task. The regressors used are: Linear Regression, Ridge Regression, and ElasticNet Regression.

```
[52]: models = {
    "Linear Regression": LinearRegression(),
    "Ridge Regression": Ridge(alpha=0.5),
    "ElasticNet Regression": ElasticNet(alpha=1.0, l1_ratio=0.5),
}

results = {}

for name, model in models.items():
    model.fit(X_train_filtered[important_genes], y_train)
    y_pred = model.predict(X_test[important_genes])

    mse = mean_squared_error(y_test, y_pred)
    r2 = r2_score(y_test, y_pred)
    results[name] = {
        "MSE": mse,
        "R2": r2
    }

for model_name, metrics in results.items():
```



```
print(f"{model_name}:\n\tMSE: {metrics['MSE']}\n\tR2: {metrics['R2']}\n")
```

Linear Regression:

```
MSE: 109.75763313259237
R2: 0.14127465204431466
```

Ridge Regression:

```
MSE: 109.66346309043713
R2: 0.14201142268986122
```

ElasticNet Regression:

```
MSE: 98.05173665922788
R2: 0.23285962645861957
```

Let's split the dataset based on sex, repeat the regression task, and compare the results between the groups.

2.0.1 Female

```
[27]: female_df_deg = df_deg[df_deg['SEX'] == 1]
```

```
[28]: female_df_deg.shape
```

```
[28]: (154, 17399)
```

```
[29]: female_meta_data = meta_data[meta_data['SEX'] == 'Female']
```

```
[30]: female_meta_data.shape
```

```
[30]: (154, 7)
```

Let's split our data and labels to train and test sets

```
[31]: X_train, X_test, y_train, y_test = train_test_split(female_df_deg,
    ↪ female_meta_data, test_size=0.2, random_state = 10101)
```

First, we applied a variance threshold to filter out features that do not have sufficient variation.

```
[32]: threshold = 0.8
selector = VarianceThreshold(threshold)
filtered_data = selector.fit_transform(X_train)
selected_features = X_train.columns[selector.get_support()]
X_train_filtered = pd.DataFrame(filtered_data, columns=selected_features)
```

```
[33]: X_train_filtered.shape
```

```
[33]: (123, 5223)
```

Next, we performed LASSO to extract the most relevant features for this task.

```
[33]: lasso = LassoCV(cv=5).fit(X_train_filtered, y_train['Age'])
# Extract the non-zero coefficients
female_important_genes = X_train_filtered.columns[(lasso.coef_ != 0)]
```

```
[34]: female_important_genes
```

```
[34]: Index(['ENSG00000007001', 'ENSG00000070193', 'ENSG00000086548',
            'ENSG00000100867', 'ENSG00000104044', 'ENSG00000104435',
            'ENSG00000112852', 'ENSG00000128313', 'ENSG00000129824',
            'ENSG00000131080', 'ENSG00000132872', 'ENSG00000134121',
            'ENSG00000134640', 'ENSG00000134668', 'ENSG00000139515',
            'ENSG00000151962', 'ENSG00000161973', 'ENSG00000162391',
            'ENSG00000162426', 'ENSG00000164266', 'ENSG00000165246',
            'ENSG00000166206', 'ENSG00000169876', 'ENSG00000171747',
            'ENSG00000174226', 'ENSG00000179674', 'ENSG00000180440',
            'ENSG00000182256', 'ENSG00000182771', 'ENSG00000196844',
            'ENSG00000198885', 'ENSG00000205358', 'ENSG00000206077',
            'ENSG00000213088', 'ENSG00000221890', 'ENSG00000223609',
            'ENSG00000236882', 'ENSG00000237693', 'ENSG00000239839',
            'ENSG00000250361', 'ENSG00000251209', 'ENSG00000255798',
            'ENSG00000256618', 'ENSG00000258484'],
            dtype='object', name='gene')
```

```
[35]: len(female_important_genes)
```

```
[35]: 44
```

```
[34]: female_df_deg[female_important_genes]
```

```
[34]: gene  ENSG00000007001  ENSG00000070193  ENSG00000086548  ENSG00000100867  \
0          3.061454          0.346095          2.234674          2.830128
1          8.213882          0.863340          2.961642          3.886277
3          9.549683          0.096771          2.112367          2.341393
4          2.541464          0.085581          1.142914          1.746493
5          4.478528          0.350495          1.645951          2.773706
..          ...          ...          ...          ...
185         7.831961         -0.678662         -2.749051         -3.486017
186         7.730441         -2.342361         -5.149716          2.584994
188         7.990447         -0.632825         -4.539715          1.241645
190         4.256733         -2.002540         -1.517113          4.977028
191         4.089949         -3.173712         -1.951319         -1.299243

gene  ENSG00000104044  ENSG00000104435  ENSG00000112852  ENSG00000128313  \
0          1.956820          1.440422          0.809067          0.265925
1          2.372354          1.460581          0.863340          1.132527
3          1.503946          1.354568          1.127797          0.617603
4          0.490838          0.430717          0.161530          0.005411
```

5	1.414625	0.920810	0.435384	-0.445972		
..		
185	-1.901054	-3.486017	-2.263624	-2.263624		
186	-2.827788	-1.979791	-1.062253	-5.149716		
188	-4.539715	-2.217787	-2.954753	-0.839275		
190	-2.002540	-2.739505	-1.517113	-0.624028		
191	-4.758674	-0.851784	-1.588749	0.667591		
gene	ENSG00000129824	ENSG00000131080	...	ENSG00000221890	ENSG00000223609	\
0	-3.278396	2.758693	...	1.706497	0.181036	
1	-2.837100	2.890821	...	1.952977	0.787391	
3	-4.426791	2.314676	...	1.913059	-0.726352	
4	-5.123872	2.555608	...	0.233680	-0.265891	
5	-4.693899	3.132649	...	1.681140	-0.606437	
..	
185	-5.070979	-5.070979	...	-0.316092	-5.070979	
186	-3.564753	-0.901788	...	0.059737	-3.564753	
188	2.201752	0.952138	...	-1.732360	-4.539715	
190	-4.324468	-0.865036	...	0.067849	-2.002540	
191	-3.173712	1.391073	...	-0.671211	-3.173712	
gene	ENSG00000236882	ENSG00000237693	ENSG00000239839	ENSG00000250361	\	
0	3.960009	1.245166	3.000828	1.019285		
1	7.135880	2.020881	-0.334599	0.935490		
3	5.746886	1.127797	-0.339329	-0.339329		
4	5.103744	0.302393	-0.265891	-0.600310		
5	6.590925	-0.050043	-2.371971	0.060988		
..		
185	4.631193	-5.070979	-0.427123	-2.749051		
186	7.257286	-3.564753	-1.979791	-2.827788		
188	6.416661	-0.632825	1.903228	-2.954753		
190	6.016495	-4.324468	1.942319	-2.739505		
191	6.967970	-4.758674	-0.235112	-4.758674		
gene	ENSG00000251209	ENSG00000255798	ENSG00000256618	ENSG00000258484		
0	0.751351	2.007006	6.124616	-0.219502		
1	1.460581	1.600305	6.772079	-0.721623		
3	0.431190	-0.967360	6.704422	-0.339329		
4	0.657488	-0.875945	5.328369	0.430717		
5	1.188744	2.486010	6.479153	-1.886544		
..		
185	-0.427123	2.148189	8.637566	-1.901054		
186	-2.827788	-1.062253	9.595696	-1.979791		
188	0.504679	-0.839275	9.543349	-0.839275		
190	-0.865036	-0.417577	9.142755	-2.739505		
191	-1.588749	-4.758674	10.316011	-1.058235		

[154 rows x 44 columns]

It's time for regression. I have used three types of regressors in this part to ensure consistency. These regressors are well-established in machine learning and are well-suited for this task. The regressors used are: Linear Regression, Ridge Regression, and ElasticNet Regression.

```
[56]: models = {
        "Linear Regression": LinearRegression(),
        "Ridge Regression": Ridge(alpha=1.0),
        "ElasticNet Regression": ElasticNet(alpha=1.0, l1_ratio=0.5),
    }

    results = {}

    for name, model in models.items():
        model.fit(X_train_filtered[female_important_genes], y_train['Age'])
        y_pred = model.predict(X_test[female_important_genes])

        mse = mean_squared_error(y_test['Age'], y_pred)
        r2 = r2_score(y_test['Age'], y_pred)
        results[name] = {
            "MSE": mse,
            "R2": r2
        }

    for model_name, metrics in results.items():
        print(f"{model_name}:\n\tMSE: {metrics['MSE']}\n\tR2: {metrics['R2]}\n")
```

Linear Regression:

MSE: 88.31876196989404
R²: 0.3969937886987882

Ridge Regression:

MSE: 88.56495689237829
R²: 0.39531286536904964

ElasticNet Regression:

MSE: 97.7445429895872
R²: 0.3326382160609207

2.0.2 Male

```
[34]: male_df_deg = df_deg[df_deg['SEX'] == 0]
```

```
[35]: male_df_deg.shape
```

```
[35]: (38, 17399)
```

```
[36]: male_meta_data = meta_data[meta_data['SEX'] == 'Male']
```

```
[37]: male_meta_data.shape
```

```
[37]: (38, 7)
```

Let's split our data and labels to train and test sets

```
[38]: X_train, X_test, y_train, y_test = train_test_split(male_df_deg,
↳ male_meta_data, test_size=0.2, random_state = 10101)
```

First, we applied a variance threshold to filter out features that do not have sufficient variation.

```
[39]: threshold = 0.8
selector = VarianceThreshold(threshold)
filtered_data = selector.fit_transform(X_train)
selected_features = X_train.columns[selector.get_support()]
X_train_filtered = pd.DataFrame(filtered_data, columns=selected_features)
```

```
[40]: X_train_filtered.shape
```

```
[40]: (30, 5016)
```

Next, we performed LASSO to extract the most relevant features for this task.

```
[45]: lasso = LassoCV(cv=5).fit(X_train_filtered, y_train['Age'])
# Extract the non-zero coefficients
male_important_genes = X_train_filtered.columns[(lasso.coef_ != 0)]
```

```
[42]: male_important_genes
```

```
[42]: Index(['ENSG000000091138', 'ENSG000000095596', 'ENSG00000106236',
'ENSG00000122121', 'ENSG00000131080', 'ENSG00000134184',
'ENSG00000140379', 'ENSG00000141668', 'ENSG00000147257',
'ENSG00000152726', 'ENSG00000168398', 'ENSG00000184674',
'ENSG00000198074', 'ENSG00000215700', 'ENSG00000235098',
'ENSG00000235961', 'ENSG00000244694'],
dtype='object', name='gene')
```

```
[43]: len(male_important_genes)
```

```
[43]: 17
```

```
[44]: X_test[male_important_genes].shape
```

```
[44]: (8, 17)
```

It's time for regression. I have used three types of regressors in this part to ensure consistency. These regressors are well-established in machine learning and are well-suited for this task. The regressors used are: Linear Regression, Ridge Regression, and ElasticNet Regression.

```
[48]: models = {
    "Linear Regression": LinearRegression(),
    "Ridge Regression": Ridge(alpha=1.0),
    "ElasticNet Regression": ElasticNet(alpha=1.0, l1_ratio=0.5),
}

results = {}

for name, model in models.items():
    model.fit(X_train_filtered[male_important_genes], y_train['Age'])
    y_pred = model.predict(X_test[male_important_genes])

    mse = mean_squared_error(y_test['Age'], y_pred)
    r2 = r2_score(y_test['Age'], y_pred)
    results[name] = {
        "MSE": mse,
        "R2": r2
    }

for model_name, metrics in results.items():
    print(f"{model_name}:\n\tMSE: {metrics['MSE']}\n\tR2: {metrics['R2]}\n")
```

Linear Regression:

MSE: 120.18807385165366
R²: -0.6186945973286688

Ridge Regression:

MSE: 121.76068452045129
R²: -0.639874538995977

ElasticNet Regression:

MSE: 132.0734434710684
R²: -0.7787669154352648

So far, we have obtained only a single result for each regressor, but these results are not statistically valid. To test for reproducibility, we have put the entire process into a loop to run it 10 times.

```
[61]: # Load Data
normal_counts = pd.read_csv('Normal.counts.voom.csv')
meta_data = pd.read_csv('meta_data.csv')

categorical_features = ['SEX', 'Diabet']
ohe = OneHotEncoder()
numerical_features = ['BMI_surg', 'Age']
scaler = StandardScaler()
preprocessor = ColumnTransformer(
    transformers=[
```

```

        ('num', scaler, numerical_features),
        ('cat', ohe, categorical_features)])
clinical_data = preprocessor.fit_transform(meta_data)
clinical_data = pd.DataFrame(clinical_data).iloc[:, [0, 2, 4]]
clinical_data.columns = ['BMI_surg', 'Sex', 'Diabet']
gene_names = normal_counts.iloc[:,0]
gene_names.at[17396] = 'BMI_surg'
gene_names.at[17397] = 'SEX'
gene_names.at[17398] = 'Diabet'

combined_data = np.hstack((normal_counts.iloc[0:,1:].T, clinical_data))
df_deg = pd.DataFrame(combined_data)
df_deg.columns = gene_names

LR_R2s = []
LR_MSEs = []
f_LR_R2s = []
f_LR_MSEs = []
m_LR_R2s = []
m_LR_MSEs = []

RR_R2s = []
RR_MSEs = []
f_RR_R2s = []
f_RR_MSEs = []
m_RR_R2s = []
m_RR_MSEs = []

ENR_R2s = []
ENR_MSEs = []
f_ENR_R2s = []
f_ENR_MSEs = []
m_ENR_R2s = []
m_ENR_MSEs = []

for i in range(10):
    print('iteration',i)
    X_train, X_test, y_train, y_test = train_test_split(df_deg,
↳ meta_data['Age'], test_size=0.2, random_state = i)
    threshold = 0.8
    selector = VarianceThreshold(threshold)
    filtered_data = selector.fit_transform(X_train)
    selected_features = X_train.columns[selector.get_support()]
    X_train_filtered = pd.DataFrame(filtered_data, columns=selected_features)
    lasso = LassoCV(cv=5).fit(X_train_filtered, y_train)
    # Extract the non-zero coefficients
    important_genes = X_train_filtered.columns[(lasso.coef_ != 0)]

```

```

# Linear Regression
LR_model = LinearRegression()
LR_model.fit(X_train_filtered[important_genes], y_train)
y_pred_LR = LR_model.predict(X_test[important_genes])
LR_MSE = mean_squared_error(y_test, y_pred_LR)
LR_R2 = r2_score(y_test, y_pred_LR)
LR_MSEs.append(LR_MSE)
LR_R2s.append(LR_R2)

# Ridge Regression
RR_model = Ridge(alpha=0.5)
RR_model.fit(X_train_filtered[important_genes], y_train)
y_pred_RR = RR_model.predict(X_test[important_genes])
RR_MSE = mean_squared_error(y_test, y_pred_RR)
RR_R2 = r2_score(y_test, y_pred_RR)
RR_MSEs.append(RR_MSE)
RR_R2s.append(RR_R2)

# ElasticNet Regression
ENR_model = ElasticNet(alpha=1.0, l1_ratio=0.5)
ENR_model.fit(X_train_filtered[important_genes], y_train)
y_pred_ENR = ENR_model.predict(X_test[important_genes])
ENR_MSE = mean_squared_error(y_test, y_pred_ENR)
ENR_R2 = r2_score(y_test, y_pred_ENR)
ENR_MSEs.append(ENR_MSE)
ENR_R2s.append(ENR_R2)

female_df_deg = df_deg[df_deg['SEX'] == 1]
female_meta_data = meta_data[meta_data['SEX'] == 'Female']
X_train, X_test, y_train, y_test = train_test_split(female_df_deg,
female_meta_data, test_size=0.2, random_state = i)
threshold = 0.8
selector = VarianceThreshold(threshold)
filtered_data = selector.fit_transform(X_train)
selected_features = X_train.columns[selector.get_support()]
X_train_filtered = pd.DataFrame(filtered_data, columns=selected_features)
lasso = LassoCV(cv=5).fit(X_train_filtered, y_train['Age'])
# Extract the non-zero coefficients
female_important_genes = X_train_filtered.columns[(lasso.coef_ != 0)]

# Linear Regression
LR_model = LinearRegression()
LR_model.fit(X_train_filtered[female_important_genes], y_train['Age'])
y_pred_LR = LR_model.predict(X_test[female_important_genes])
LR_MSE = mean_squared_error(y_test['Age'], y_pred_LR)

```



```

LR_R2 = r2_score(y_test['Age'], y_pred_LR)
f_LR_MSEs.append(LR_MSE)
f_LR_R2s.append(LR_R2)

# Ridge Regression
RR_model = Ridge(alpha=0.5)
RR_model.fit(X_train_filtered[female_important_genes], y_train['Age'])
y_pred_RR = RR_model.predict(X_test[female_important_genes])
RR_MSE = mean_squared_error(y_test['Age'], y_pred_RR)
RR_R2 = r2_score(y_test['Age'], y_pred_RR)
f_RR_MSEs.append(RR_MSE)
f_RR_R2s.append(RR_R2)

# ElasticNet Regression
ENR_model = ElasticNet(alpha=1.0, l1_ratio=0.5)
ENR_model.fit(X_train_filtered[female_important_genes], y_train['Age'])
y_pred_ENR = ENR_model.predict(X_test[female_important_genes])
ENR_MSE = mean_squared_error(y_test['Age'], y_pred_ENR)
ENR_R2 = r2_score(y_test['Age'], y_pred_ENR)
f_ENR_MSEs.append(ENR_MSE)
f_ENR_R2s.append(ENR_R2)

male_df_deg = df_deg[df_deg['SEX'] == 0]
male_meta_data = meta_data[meta_data['SEX'] == 'Male']
X_train, X_test, y_train, y_test = train_test_split(male_df_deg,
male_meta_data, test_size=0.2, random_state = i)
threshold = 0.8
selector = VarianceThreshold(threshold)
filtered_data = selector.fit_transform(X_train)
selected_features = X_train.columns[selector.get_support()]
X_train_filtered = pd.DataFrame(filtered_data, columns=selected_features)
lasso = LassoCV(cv=5).fit(X_train_filtered, y_train['Age'])
# Extract the non-zero coefficients
male_important_genes = X_train_filtered.columns[(lasso.coef_ != 0)]

# Linear Regression
LR_model = LinearRegression()
LR_model.fit(X_train_filtered[male_important_genes], y_train['Age'])
y_pred_LR = LR_model.predict(X_test[male_important_genes])
LR_MSE = mean_squared_error(y_test['Age'], y_pred_LR)
LR_R2 = r2_score(y_test['Age'], y_pred_LR)
m_LR_MSEs.append(LR_MSE)
m_LR_R2s.append(LR_R2)

# Ridge Regression
RR_model = Ridge(alpha=0.5)

```

```

RR_model.fit(X_train_filtered[male_important_genes], y_train['Age'])
y_pred_RR = RR_model.predict(X_test[male_important_genes])
RR_MSE = mean_squared_error(y_test['Age'], y_pred_RR)
RR_R2 = r2_score(y_test['Age'], y_pred_RR)
m_RR_MSEs.append(RR_MSE)
m_RR_R2s.append(RR_R2)

# ElasticNet Regression
ENR_model = ElasticNet(alpha=1.0, l1_ratio=0.5)
ENR_model.fit(X_train_filtered[male_important_genes], y_train['Age'])
y_pred_ENR = ENR_model.predict(X_test[male_important_genes])
ENR_MSE = mean_squared_error(y_test['Age'], y_pred_ENR)
ENR_R2 = r2_score(y_test['Age'], y_pred_ENR)
m_ENR_MSEs.append(ENR_MSE)
m_ENR_R2s.append(ENR_R2)

```

```

iteration 0
iteration 1
iteration 2
iteration 3
iteration 4
iteration 5
iteration 6
iteration 7
iteration 8
iteration 9

```

Last but not least, we need to report the average and confidence interval for the results of each classifier. We will also create a bar plot to visually compare the results of the classifiers. We also need to compare the results of the three conditions to determine whether splitting the data based on sex will affect the performance or not

```

[70]: def evaluation(LR_MSEs, LR_R2s, RR_MSEs, RR_R2s, ENR_MSEs, ENR_R2s):

    LR_mean_MSE = np.mean(LR_MSEs)
    LR_mean_R2 = np.mean(LR_R2s)
    LR_MSE_conf_interval = np.percentile(LR_MSEs, [2.5, 97.5])
    LR_R2_conf_interval = np.percentile(LR_R2s, [2.5, 97.5])
    print(f'Mean MSE for Linear Regression: {LR_mean_MSE}, 95% CI:␣
↪{LR_MSE_conf_interval}')
    print(f'Mean R2 for Linear Regression: {LR_mean_R2}, 95% CI:␣
↪{LR_R2_conf_interval}')
    ␣
    ↪print("-----

    RR_mean_MSE = np.mean(RR_MSEs)
    RR_mean_R2 = np.mean(RR_R2s)
    RR_MSE_conf_interval = np.percentile(RR_MSEs, [2.5, 97.5])

```

```

RR_R2_conf_interval = np.percentile(RR_R2s, [2.5, 97.5])
print(f'Mean MSE for Ridge Regression: {RR_mean_MSE}, 95% CI:␣
↳{RR_MSE_conf_interval}')
print(f'Mean R2 for Ridge Regression: {RR_mean_R2}, 95% CI:␣
↳{RR_R2_conf_interval}')
␣
↳print("-----

ENR_mean_MSE = np.mean(ENR_MSEs)
ENR_mean_R2 = np.mean(ENR_R2s)
ENR_MSE_conf_interval = np.percentile(ENR_MSEs, [2.5, 97.5])
ENR_R2_conf_interval = np.percentile(ENR_R2s, [2.5, 97.5])
print(f'Mean MSE for ElasticNet Regression: {ENR_mean_MSE}, 95% CI:␣
↳{ENR_MSE_conf_interval}')
print(f'Mean R2 for ElasticNet Regression: {ENR_mean_R2}, 95% CI:␣
↳{ENR_R2_conf_interval}')

model_results = {
    'Linear Regression': {
        'mean_MSE': LR_mean_MSE,
        'MSE_ci': LR_MSE_conf_interval,
        'mean_R2': LR_mean_R2,
        'R2_ci': LR_R2_conf_interval
    },
    'Ridge Regression': {
        'mean_MSE': RR_mean_MSE,
        'MSE_ci': RR_MSE_conf_interval,
        'mean_R2': RR_mean_R2,
        'R2_ci': RR_R2_conf_interval
    },
    'ElasticNet Regression': {
        'mean_MSE': ENR_mean_MSE,
        'MSE_ci': ENR_MSE_conf_interval,
        'mean_R2': ENR_mean_R2,
        'R2_ci': ENR_R2_conf_interval
    }
}

models = list(model_results.keys())
mean_MSEs = [model_results[model]['mean_MSE'] for model in models]
MSE_cis = [model_results[model]['MSE_ci'] for model in models]
mean_R2s = [model_results[model]['mean_R2'] for model in models]
R2_cis = [model_results[model]['R2_ci'] for model in models]
MSE_errors = np.array([[mean - ci[0], ci[1] - mean] for mean, ci in␣
↳zip(mean_MSEs, MSE_cis)]).T

```

```

R2_errors = np.array([[mean - ci[0], ci[1] - mean] for mean, ci in
↪zip(mean_R2s, R2_cis)]).T

fig, ax = plt.subplots(1, 2, figsize=(14, 6))

# Plot MSE
ax[0].bar(models, mean_MSEs, yerr=MSE_errors, capsize=5, color='skyblue')
ax[0].set_title('Mean MSE with 95% Confidence Interval')
ax[0].set_ylabel('MSE')
ax[0].set_ylim([20, 200])
ax[0].set_xticklabels(models, rotation=45, ha="right")

# Plot R2
ax[1].bar(models, mean_R2s, yerr=R2_errors, capsize=5, color='lightgreen')
ax[1].set_title('Mean R2 with 95% Confidence Interval')
ax[1].set_ylabel('R2')
ax[1].set_ylim([0, 1])
ax[1].set_xticklabels(models, rotation=45, ha="right")

plt.tight_layout()
plt.show()

```

1- Men and women

[71]: evaluation(LR_MSEs,LR_R2s,RR_MSEs,RR_R2s,ENR_MSEs,ENR_R2s)

Mean MSE for Linear Regression: 93.72387171724446, 95% CI: [76.13892279
122.3389011]

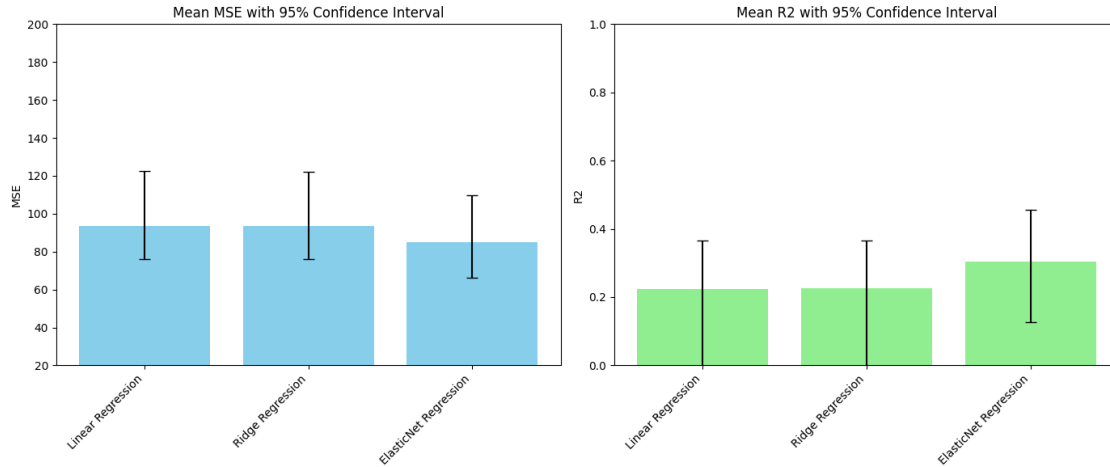
Mean R2 for Linear Regression: 0.22398519282954946, 95% CI: [-0.15693711
0.3652658]

Mean MSE for Ridge Regression: 93.58325500454752, 95% CI: [75.99978915
122.15512382]

Mean R2 for Ridge Regression: 0.2252074305987804, 95% CI: [-0.15420639
0.36603938]

Mean MSE for ElasticNet Regression: 85.03072068605735, 95% CI: [66.24003299
109.73508623]

Mean R2 for ElasticNet Regression: 0.30500370984307895, 95% CI: [0.12720519
0.4558891]



2- Women only

```
[72]: evaluation(f_LR_MSEs,f_LR_R2s,f_RR_MSEs,f_RR_R2s,f_ENR_MSEs,f_ENR_R2s)
```

Mean MSE for Linear Regression: 81.4039767743786, 95% CI: [61.99474362 102.30060272]


Mean R2 for Linear Regression: 0.32655859857434866, 95% CI: [0.07109799 0.52107335]

Mean MSE for Ridge Regression: 81.26961196554569, 95% CI: [61.87980603 102.27423771]

Mean R2 for Ridge Regression: 0.3277495253969612, 95% CI: [0.07221842 0.52102313]

Mean MSE for ElasticNet Regression: 82.14879724899781, 95% CI: [65.95899767 102.71718425]

Mean R2 for ElasticNet Regression: 0.3258626266544995, 95% CI: [0.14689447 0.45795909]



3- Men only

```
[73]: evaluation(m_LR_MSEs,m_LR_R2s,m_RR_MSEs,m_RR_R2s,m_ENR_MSEs,m_ENR_R2s)
```

Mean MSE for Linear Regression: 134.5769243894009, 95% CI: [77.24690153 277.44700226]


Mean R2 for Linear Regression: -0.4490993665624063, 95% CI: [-1.73943844 0.3035828]

Mean MSE for Ridge Regression: 124.60972408899434, 95% CI: [75.26276249 244.97374347]

Mean R2 for Ridge Regression: -0.33927324390218344, 95% CI: [-1.46493756 0.30066489]

Mean MSE for ElasticNet Regression: 115.57231466443727, 95% CI: [75.20562135 168.87937051]

Mean R2 for ElasticNet Regression: -0.25273830383631973, 95% CI: [-1.23332149 0.28383753]



As you can see, we were unable to achieve a satisfactory R^2 for any of the conditions. While it is biologically and computationally feasible to estimate human age based on gene expression data, our results suggest that the limited sample size was a significant factor in our lack of success. With fewer than 200 samples, there was insufficient data to effectively perform this regression task. Additionally, although we hypothesized that splitting the data based on sex might improve age prediction, this approach further reduced the number of samples (to fewer than 40 for men), making it impossible for us to accomplish our desired results.